

# Yao Dou

douy@gatech.edu

<https://yao-dou.github.io/>

## Education

- 2021 Georgia Institute of Technology  
Ph.D. in Computer Science Aug 2021 – Aug 2026  
Advisor: Prof. Wei Xu
- 2017 University of Washington  
B.S. in Computer Science Sep 2017 – Jun 2021  
Graduated at age 20  
Advisor: Prof. Yejin Choi

## Experience

- 2024 Microsoft Research, Deep Learning Group  
Research Intern, May 2024 – April 2025  
Mentors: Michel Galley and Jianfeng Gao
- 2021 Allen Institute for AI, Mosaic Team  
Research Intern, Jan 2021 – Mar 2021  
Mentors: Maxwell Forbes and Yejin Choi
- 2020 Noonum, Inc  
Data Science Intern, Jun 2020 – Sep 2020

## Publications

- EMNLP 2025 SimulatorArena: Are User Simulators Reliable Proxies for Multi-Turn Evaluation of AI Assistants?  
**Yao Dou**, Michel Galley, Baolin Peng, Chris Kedzie, Weixin Cai, Alan Ritter, Chris Quirk, Wei Xu, Jianfeng Gao
- COLM 2025 Evaluating LLMs on Chinese Idiom Translation  
Cai Yang, **Yao Dou**, David Heineman, Xiaofeng Wu, Wei Xu
- ICML 2025 CollabLLM: From Passive Responders to Active Collaborators  
Shirley Wu, Michel Galley, Baolin Peng, Hao Cheng, Gavin Li, **Yao Dou**, Weixin Cai, James Zou, Jure Leskovec, Jianfeng Gao

## Outstanding Paper Award

- CSCW 2025 Measuring, Modeling, and Helping People Account for Privacy Risks in Online Self-Disclosures with AI  
Isadora Krsek, Anubha Kabra, **Yao Dou**, Tarek Namus, Laura A Dabbish, Alan Ritter, Wei Xu, Sauvik Das
- AAAI 2025 CROSSNEWS: A Cross-Genre Authorship Verification and Attribution Benchmark  
Marcus Ma, Duong Minh Le, Junmo Kang, **Yao Dou**, John Cadigan, Dayne Freitag, Alan Ritter, Wei Xu
- ArXiv TemporalBench: Benchmarking Fine-grained Temporal Understanding for Multimodal Video Models  
Mu Cai, Reuben Tan, Jianrui Zhang, Bocheng Zou, Kai Zhang, Feng Yao, Fangrui Zhu, Jing Gu, Yiwu Zhong, Yuzhang Shang, **Yao Dou**, Jaden Park, Jianfeng Gao, Yong Jae Lee, Jianwei Yang
- EMNLP 2024 Improving Minimum Bayes Risk Decoding with Multi-Prompt  
David Heineman, **Yao Dou**, Wei Xu
- EMNLP 2024 GPT-4 Jailbreaks Itself with Near-Perfect Success Using Self-Explanation  
Govind Ramesh, **Yao Dou**, Wei Xu
- ACL 2024 Reducing Privacy Risks in Online Self-Disclosures with Language Models  
**Yao Dou**, Isadora Krsek, Tarek Naous, Anubha Kabra, Sauvik Das, Alan Ritter, Wei Xu
- EMNLP 2023 Thresh: A Unified, Customizable and Deployable Platform for Fine-grained Text Evaluation (Demo)  
David Heineman, **Yao Dou**, Wei Xu
- EMNLP 2023 Dancing Between Success and Failure: Edit-Level Simplification Evaluation using SALSA  
David Heineman, **Yao Dou**, Mounica Maddela, Wei Xu
- ACL 2023 LENS: A learnable Evaluation Metric for Text Simplification  
Mounica Maddela\*, **Yao Dou\***, David Heineman, Wei Xu
- EMNLP 2022 Improving Large-scale Paraphrase Acquisition and Generation  
**Yao Dou**, Chao Jiang, Wei Xu
- ACL 2022 Is GPT-3 Text Indistinguishable from Human Text? Scarecrow: A Framework for Scrutinizing Machine Text  
**Yao Dou\***, Maxwell Forbes\*, Rik Koncel-Kedziorski, Noah A. Smith, Yejin Choi
- AAAI 2021 MultiTalk: A Highly-Branching Dialog Testbed for Diverse Conversations  
**Yao Dou**, Maxwell Forbes, Ari Holtzman, Yejin Choi

## Tutorials

- ACL 2024 Automatic and Human-AI Interactive Text Generation  
**Yao Dou\***, Philippe Laban\*, Claire Gardent, Wei Xu

## Honors and Awards

Apple Fellowship Finalist (1 of 3 at Georgia Tech), 2024

Microsoft Scholarship, 2020 – 2021

Dean's List, 2017 – 2021

## Teaching

Georgia Institute of Technology

TA, Advanced NLP (CS 8803), Fall 2023; NLP (CS 4650), Spring 2025

Prof. Wei Xu

TA, Machine Learning (CS 4641), Fall 2022

Prof. Mahdi Roozbahani

University of Washington

TA, Machine Learning (CSE 446), Spring 2020 and Fall 2020

Prof. Kevin Jamieson, Prof. Jamie Morgenstern

## Undergraduate Advising

David Heineman -> PYI at Allen Institute for AI

Cai Yang

Govind Ramesh

Jerry Lou Zheng

Vinayak Athavale