

Yao Dou

douy@gatech.edu

<https://yao-dou.github.io/>

Education

- 2021 Georgia Institute of Technology
Ph.D. in Computer Science 2021 – 2026
Advisor: Prof. Wei Xu
- 2017 University of Washington
B.S. in Computer Science 2017 – 2021
Magna Cum Laude
Advisor: Prof. Yejin Choi

Experience

- 2024 Microsoft Research, Deep Learning Group
Research Intern, May – Aug 2024
Mentor: Chris Kedzie and Michel Galley
- 2021 Allen Institute for AI, Mosaic Team
Research Intern, Jan 2021 – Mar 2021
- 2020 Noonum, Inc
Data Science Intern, Jun 2020 – Sep 2020

Preprints

- arXiv Jul 24 Improving Minimum Bayes Risk Decoding with Multi-Prompt
David Heineman, Yao Dou, Wei Xu
- arXiv May 24 GPT-4 Jailbreaks Itself with Near-Perfect Success Using Self-Explanation
Govind Ramesh, Yao Dou, Wei Xu

Publications

- ACL 2024 Reducing Privacy Risks in Online Self-Disclosures with Language Models
Yao Dou, Isadora Krsek, Tarek Naous, Anubha Kabra, Sauvik Das, Alan Ritter, Wei Xu
- EMNLP 2023 Thresh: A Unified, Customizable and Deployable Platform for Fine-grained Text Evaluation (*Demo*)
David Heineman, Yao Dou, Wei Xu

- EMNLP 2023 Dancing Between Success and Failure: Edit-Level Simplification Evaluation using SALSA
David Heineman, Yao Dou, Mounica Maddela, Wei Xu
- ACL 2023 LENS: A learnable Evaluation Metric for Text Simplification
Mounica Maddela, Yao Dou*, David Heineman, Wei Xu*
- EMNLP 2022 Improving Large-scale Paraphrase Acquisition and Generation
Yao Dou, Chao Jiang, Wei Xu
- ACL 2022 Is GPT-3 Text Indistinguishable from Human Text? Scarecrow: A Framework for Scrutinizing Machine Text
Yao Dou, Maxwell Forbes*, Rik Koncel-Kedziorski, Noah A. Smith, Yejin Choi*
- AAAI 2021 MultiTalk: A Highly-Branching Dialog Testbed for Diverse Conversations
Yao Dou, Maxwell Forbes, Ari Holtzman, Yejin Choi

Tutorials

- ACL 2024 Automatic and Human-AI Interactive Text Generation
Yao Dou, Philippe Laban*, Claire Gardent, Wei Xu*

Honors and Awards

Microsoft Scholarship, 2020 – 2021
Dean's List, 2017 – 2021

Teaching

Georgia Institute of Technology

TA, Advanced NLP (CS 8803), Fall 2023
Prof. Wei Xu

TA, Machine Learning (CS 4641), Fall 2022
Prof. Mahdi Roozbahani

University of Washington

TA, Machine Learning (CSE 446), Spring 2020 and Fall 2020
Prof. Kevin Jamieson, Prof. Jamie Morgenstern

Undergraduates Mentoring

David Heineman

Cai Yang

Vinayak Athavale

Govind Ramesh